

ENGR-E 511; ENGR-E 399

Machine Learning for Signal Processing

Module 07: Hidden Markov Models

Minje Kim

Department of Intelligent Systems Engineering

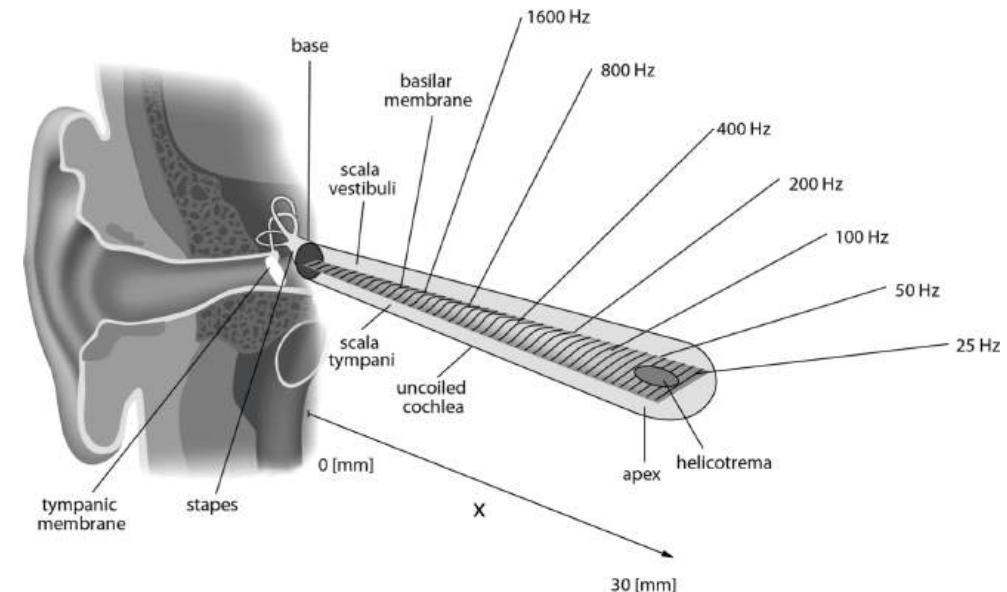
Email: minje@indiana.edu

Website: <http://minjekim.com>

Research Group: <http://saige.sice.indiana.edu>

Meeting Request: <http://doodle.com/minje>

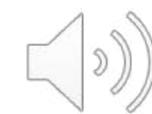
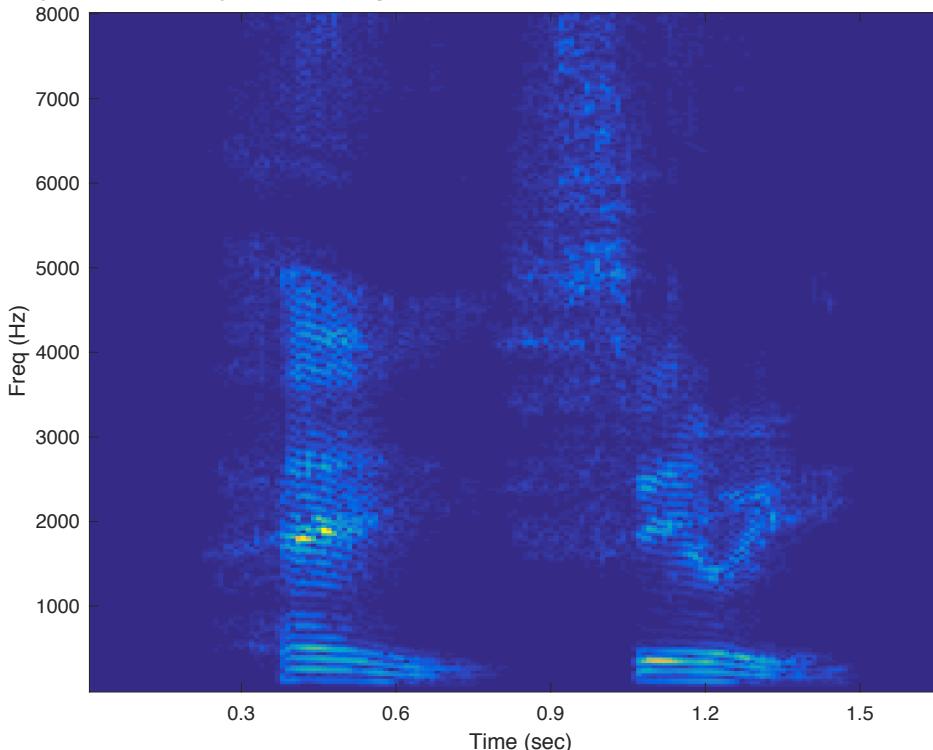
INDIANA UNIVERSITY
**SCHOOL OF INFORMATICS,
COMPUTING, AND ENGINEERING**



Keyword Spotting

- Preprocessing

- I like Siri
 - But, she doesn't like me back
- What should I do to build my own keyword spotter?



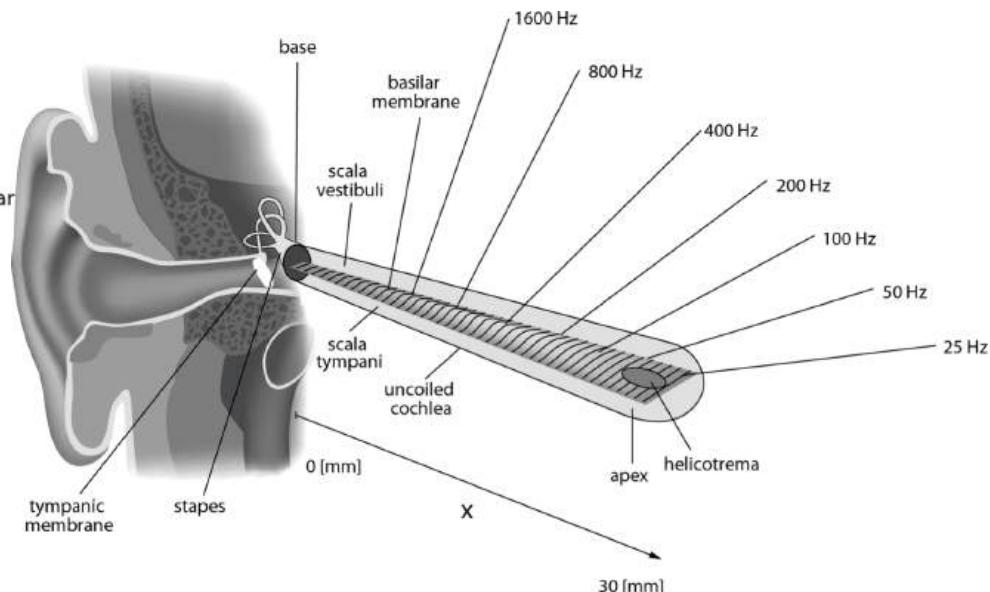
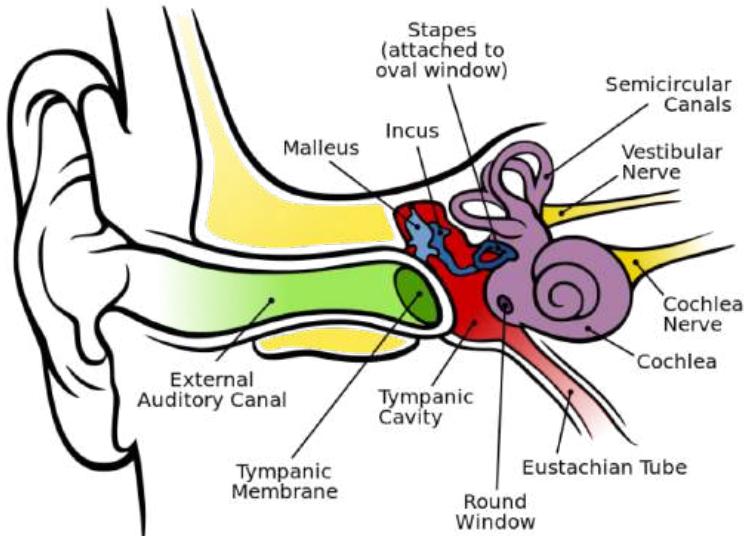
INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Keyword Spotting

- Preprocessing

- Problem with STFT
 - Too many dimensions
 - Even before that, our auditory system is not linear



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

<https://en.wikipedia.org/wiki/Electrocochleography>

<https://ocw.mit.edu/courses/health-sciences-and-technology/hst-730-molecular-biology-for-the-auditory-system-fall-2002/hst-730f02.jpg>

Keyword Spotting

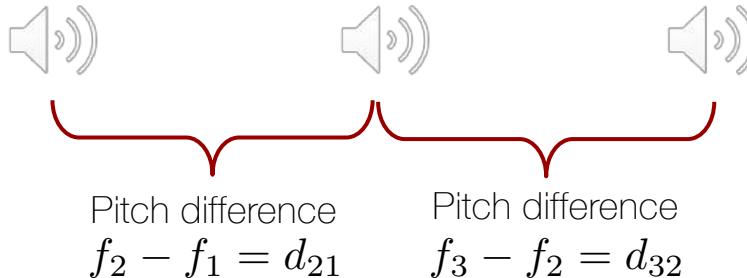
- Preprocessing

- Nonlinear frequency?

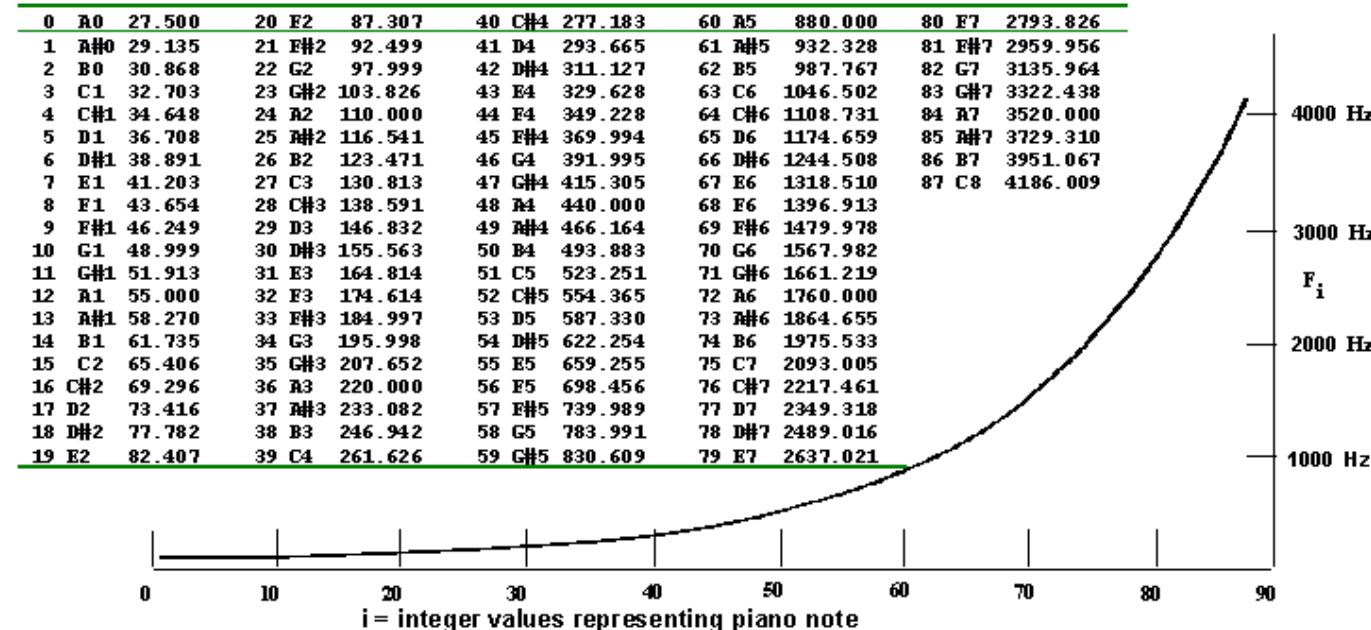
$$f_1 = 622.25$$

$$f_2 = 698.46$$

$$f_3 = 783.99$$



- $d_{21} = d_{32}$?
 - $d_{21} = 76.21 \neq d_{32} = 85.53$



Ideal frequencies for equally tempered 12-tone chromatic scale.

$$F_i = 2^{(i/12)} \cdot C$$

Where C = ideal frequency of A₀ = 27.5 Hz



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

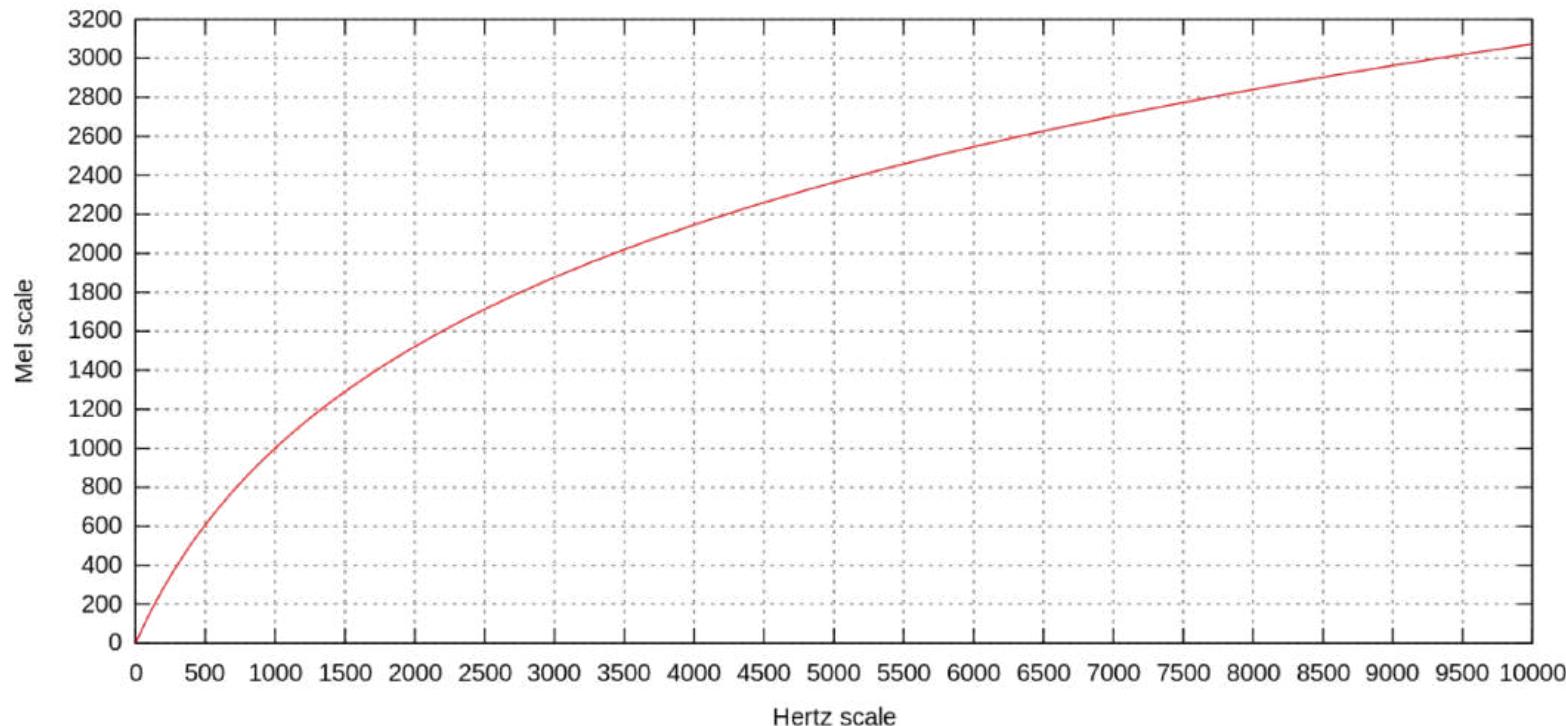
http://www.tedknowlton.com/resume/SCALE_GRAPH.gif

Keyword Spotting

- Preprocessing

- Mel scale

$$mel(f) = a \ln \left(1 + \frac{f}{b} \right) = 1127 \ln \left(1 + \frac{f}{700} \right)$$



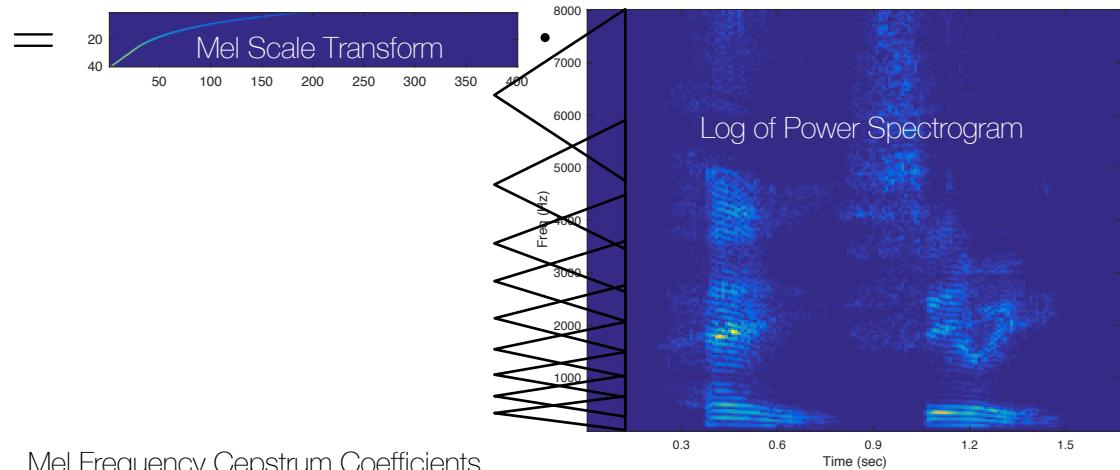
INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Keyword Spotting

- Preprocessing

- Mel scale spectra



- Still too many dimensions

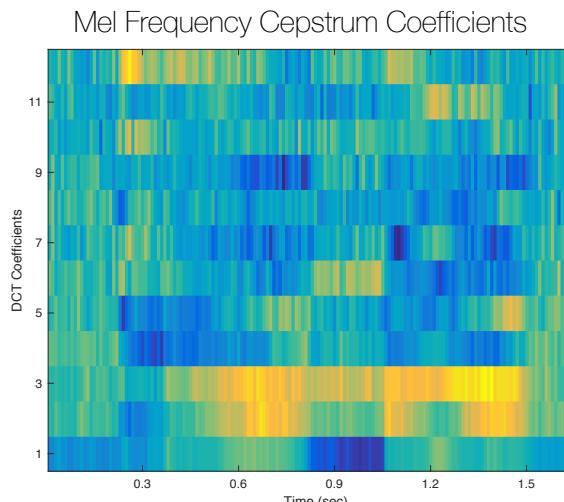
- PCA?
 - Instead, DCT

- Mel Frequency Cepstrum Coefficients (MFCC)

- From now on I will use spectrogram in the following slides to visualize the speech signal

- But, actual processing was done on the MFCC matrix

- I just don't know how to intuitively interpret MFCC

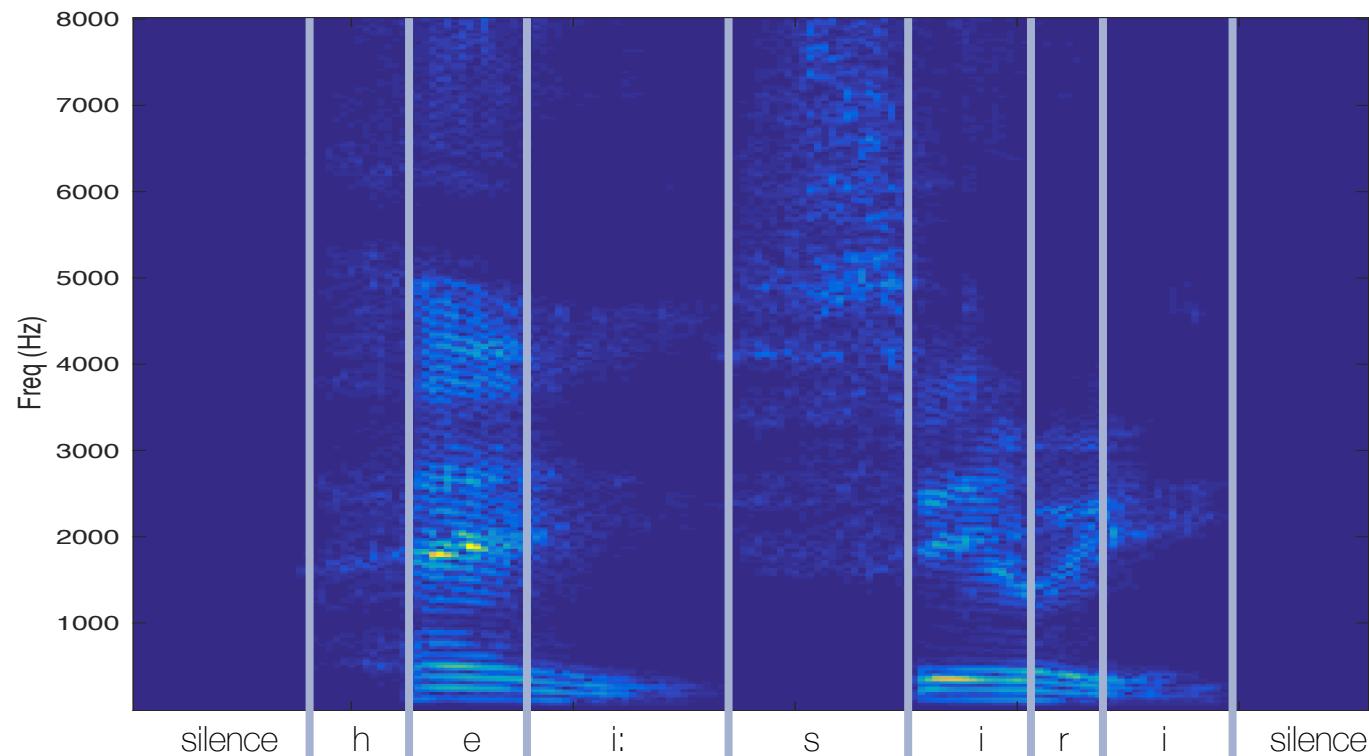


INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Training from a Sequence: a Preview

- Smoothing clustering with transition probabilities
 - I want to build a system that recognize my keyword. What should I do?
 - Prepare training signals
 - Then what?
 - I feel like learning a model out of this 12X167 data matrix.
 - What should I do?
 - Clustering!
 - But to build a model, not for grouping
 - Something similar to learning GMM for classification
 - Let me do clustering manually to build ground-truth labels



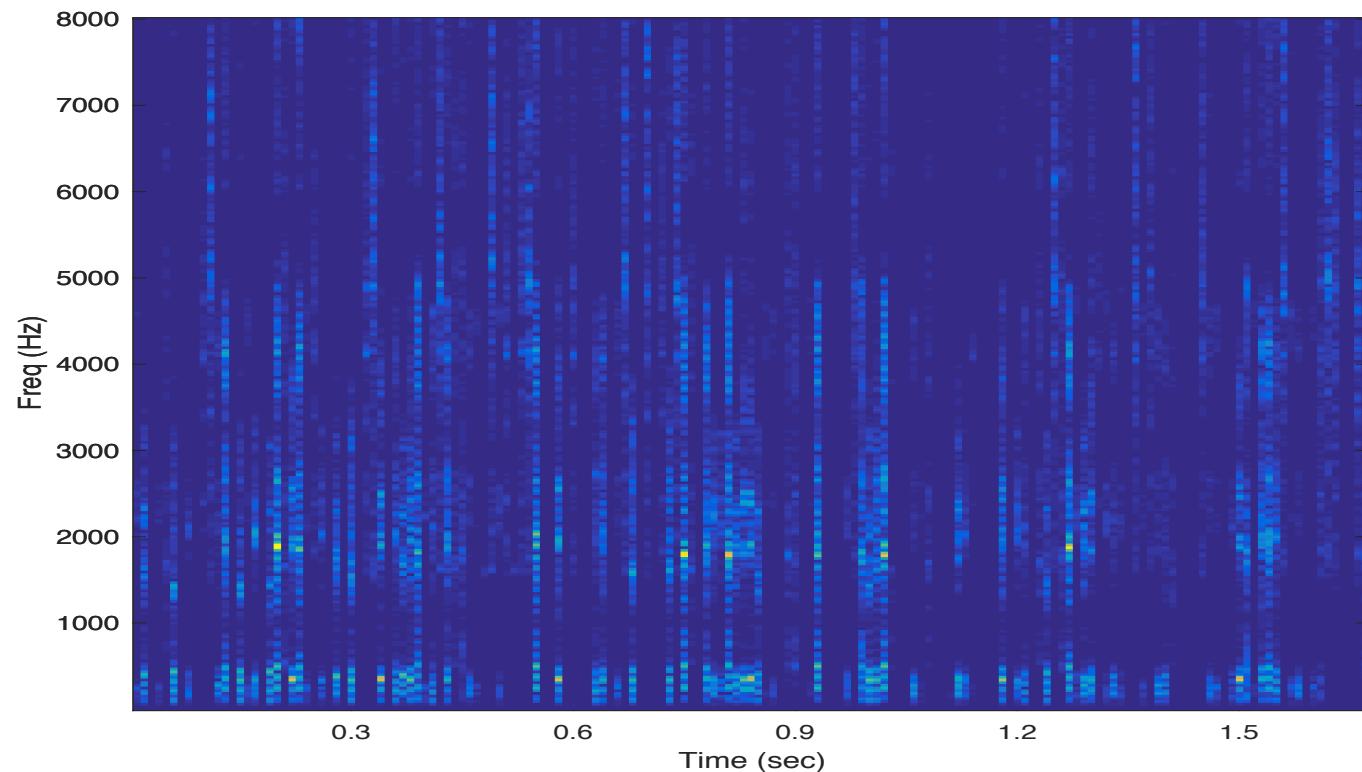
INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Training from a Sequence: a Preview

- Smoothing clustering with transition probabilities

- A big question
 - Can I do the same clustering if I scramble the column vectors?
 - Seems difficult
- Ideal clustering considers the temporal structure
 - A phoneme should sustain for a while, etc
- For GMM or Kmeans, the algorithms doesn't know the order of the data
 - Random shuffling will result in same results
 - We may need an algorithm that benefits from the temporal structure
- Will revisit this problem later in this lecture



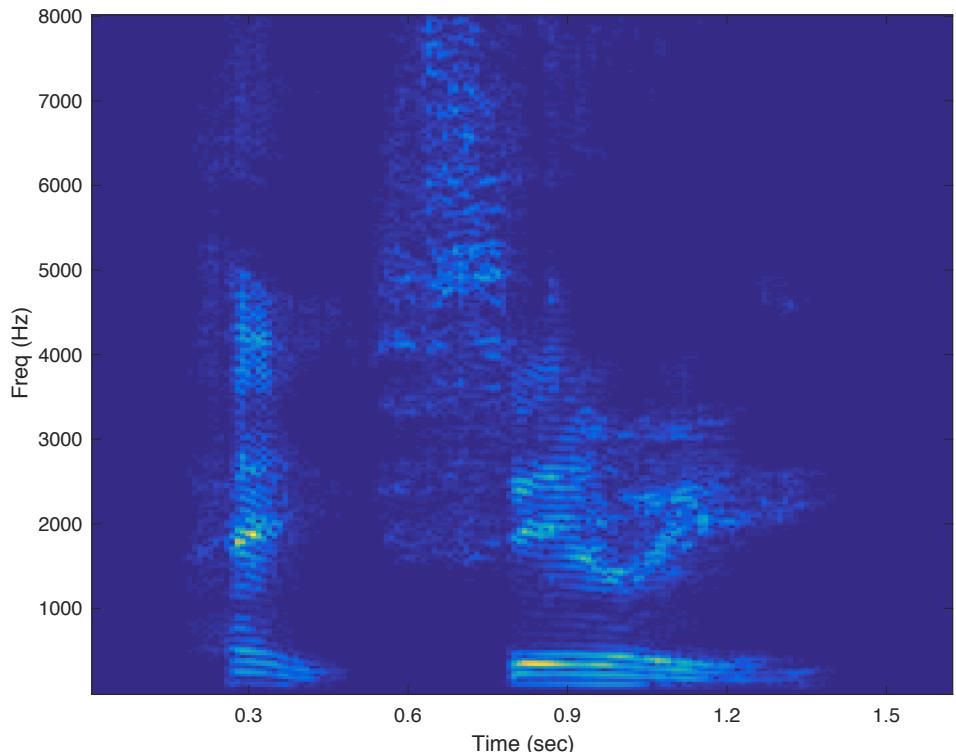
INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 1: Decoding a Sequence

- Smoothing classification with transition probabilities

- For convenience, let's assume that I know the ground-truth cluster means
 - So, we're going to assume that training is done
 - We'll revisit the training part later
- Then the problem boils down to a **series** of phoneme classification job
 - I observe a new sequence, a test signal
 - I want to assign the these new MFCC spectra to the clusters I (manually) defined before



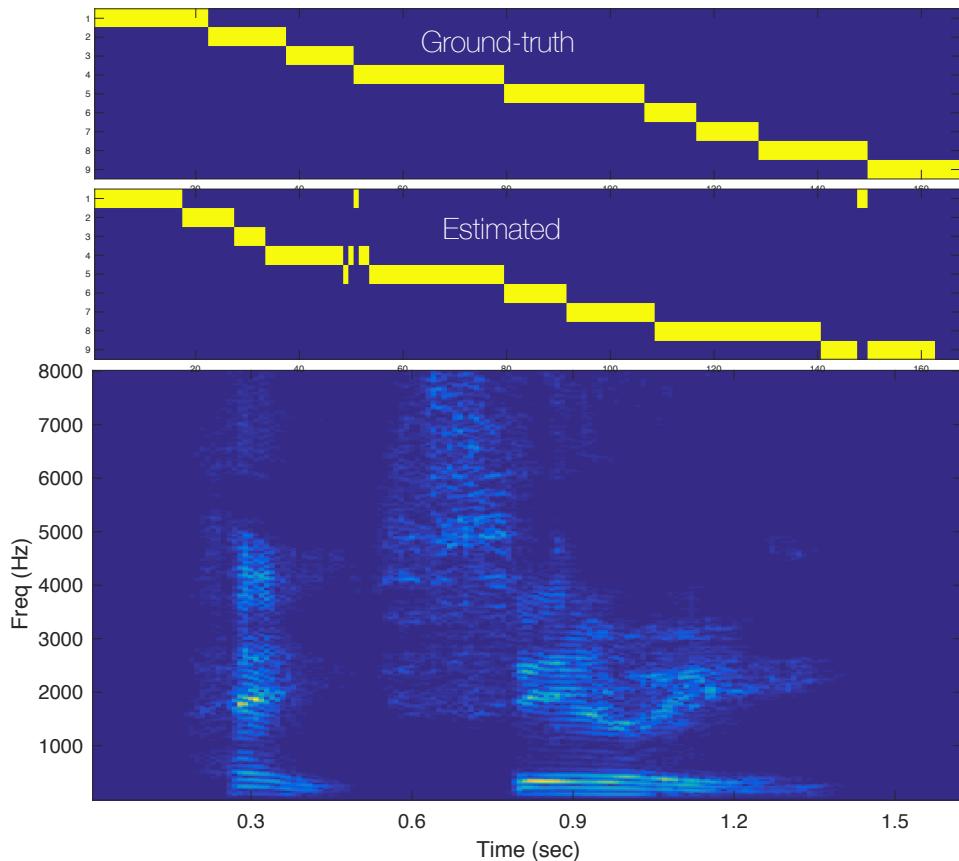
INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 1: Decoding a Sequence

- Smoothing classification with transition probabilities

- This seems like an easy naïve Bayes classification
 - As we assume known parameters
- But it's not
- This is a **decoding** problem
 - Find out the best **sequence of states** that might have generated the observed sequence
 - In our case...
 - Observation: MFCCs
 - States: Phonemes



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 1: Decoding a Sequence

- Smoothing classification with transition probabilities
 - Why is decoding difficult?
 - In a perfect world the **emission probabilities** are clean and representative

$$P(\omega_{j_t} | \mathbf{x}_t) = \frac{P(\mathbf{x}_t | \omega_{j_t}) P(\omega_{j_t})}{\sum_{\omega_{j_t}=1}^J P(\mathbf{x}_t | \omega_{j_t}) P(\omega_{j_t})}$$

- In the real world we need to deal with a test signal that can be different from the model
 - Same for all classification problem
- We tried our best as for the data generation procedure
 - i.e. we used ground-truth labels $P(\mathbf{x}_t | \omega_{j_t})$
 - But it wasn't good enough
- What can we do to improve the performance?
 - $P(\omega_{j_t}) \neq P(\omega_{j_t} | \omega_{j_{t-1}}, \omega_{j_{t-2}}, \dots, \omega_{j_1}) = P(\omega_{j_t} | \omega_{j_{t-1}})$
 - A first-order **Markov** model



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 1: Decoding a Sequence

- Smoothing classification with transition probabilities
 - The first-order Markov model for our example (in English)
 - “If the previous frame was C1 (silence), the current one would be C1 with a high probability”
$$P(\omega_{j_t} = 1 | \omega_{j_{t-1}} = 1) = 0.8$$
 - “But with a small probability, there’s a chance that the current frame can be C2 ('h')”
$$P(\omega_{j_t} = 2 | \omega_{j_{t-1}} = 1) = 0.2$$
 - “I’m pretty sure that the current frame won’t be any of the phonemes from C3 to C9”
$$P(\omega_{j_t} = 3 | \omega_{j_{t-1}} = 1) = P(\omega_{j_t} = 4 | \omega_{j_{t-1}} = 1) = \dots = P(\omega_{j_t} = 9 | \omega_{j_{t-1}} = 1) = 0$$
 - This kind of intuition can build a set of **transition probabilities**
 - It works like *a priori* distribution about the states (clusters) before even seeing the observations



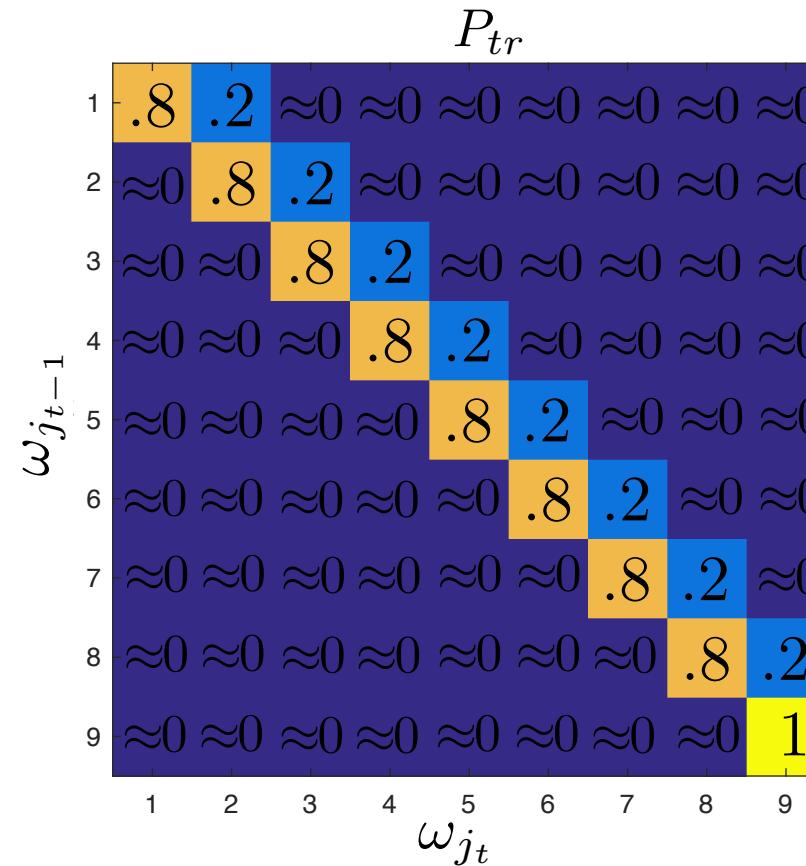
INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 1: Decoding a Sequence

- Smoothing classification with transition probabilities

- The transition probability matrix



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 1: Decoding a Sequence

- Smoothing classification with transition probabilities

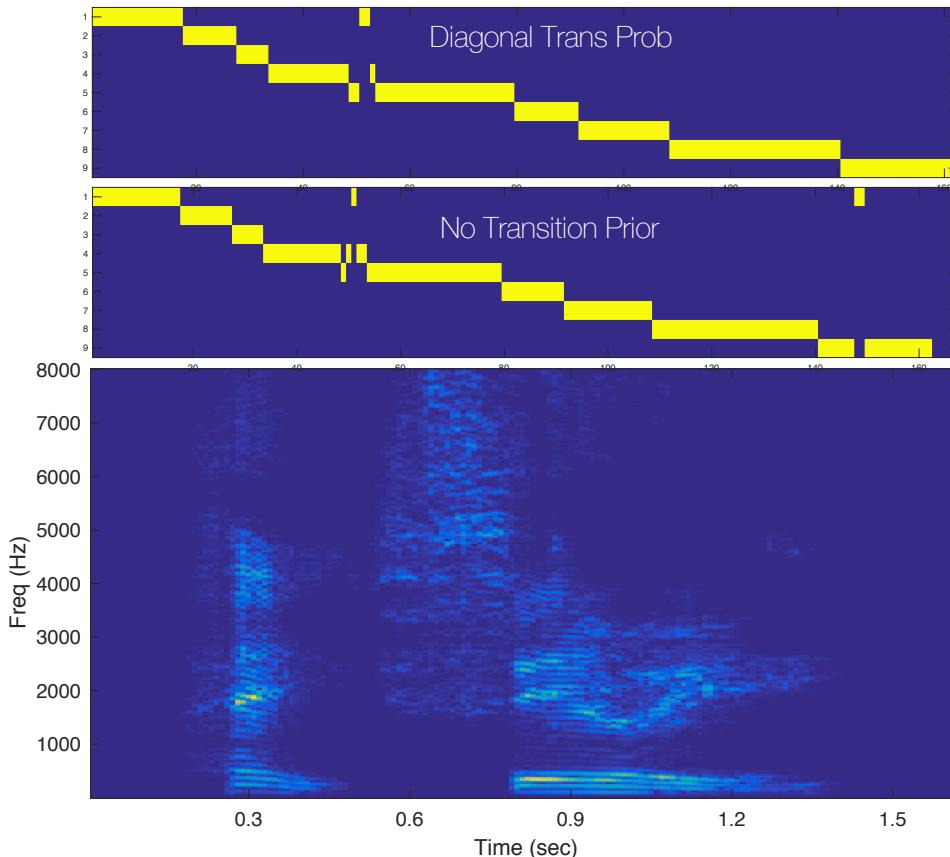
- At a given time frame I have an observation \tilde{x}_t

- The goal is to estimate the posterior probability (actually the joint prob) at every frame

$$\begin{aligned} P(\omega_{j_t} | \tilde{x}_t) &\propto P(\tilde{x}_t | \omega_{j_t}) P(\omega_{j_t} | \omega_{j_{t-1}}) \\ &= \mathcal{N}(\tilde{x}_t; \mu_{j_t}, \Sigma_{j_t}) P(\omega_{j_t} | \omega_{j_{t-1}}) \end{aligned}$$

- In English

- Even if my Gaussian distribution says that the frame should be C1,
 - I ignore it and assign it to C9 as its previous frame was C9



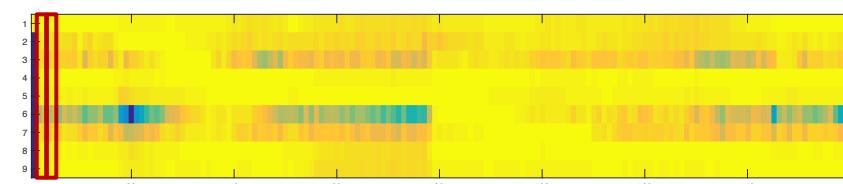
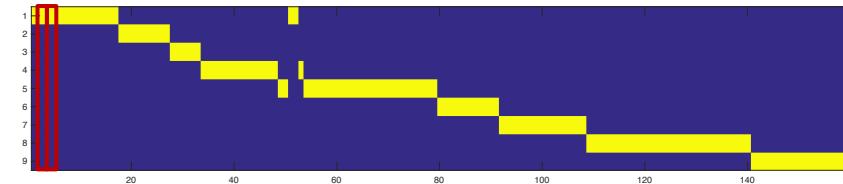
INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 1: Decoding a Sequence

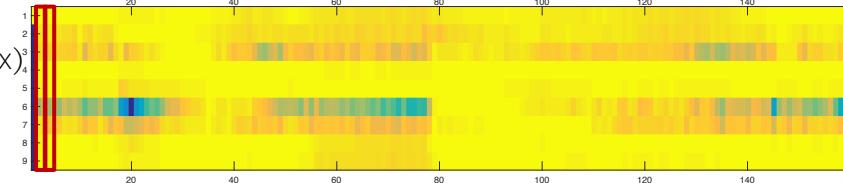
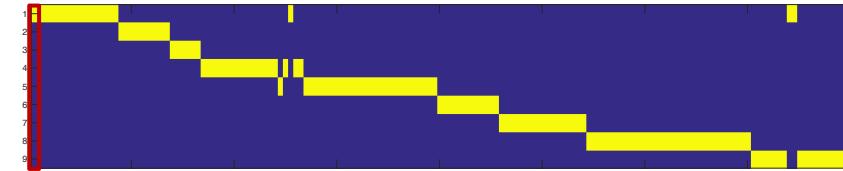
- Smoothing classification with transition probabilities

$$P_{tr}^\top$$



$$P_{tr}^\top$$

Transition matrix
(Note: no time index)



$$\tilde{U}_{(j,t)} = \begin{cases} 1 & \text{if } P_{\mathbf{U}_{(j,t)}} \odot \mathbf{U}_{(j,t)} \\ & = \max_j P_{\mathbf{U}_{(j,t)}} \odot \mathbf{U}_{(j,t)} \\ 0 & \text{otherwise} \end{cases}$$

Final labels from posterior dist

$$P_{\mathbf{U}} \odot \mathbf{U}$$

$$\log P_{\mathbf{U}} \odot \mathbf{U} = \log P_{\mathbf{U}} + \log \mathbf{U}$$

$$P_{\mathbf{U}} = P(\omega_{j_t} | \omega_{j_{t-1}}) = P_{tr}^\top \tilde{\mathbf{U}}(:, t-1)$$

$$\bar{U}_{(j,t)} = \begin{cases} 1 & \text{if } \mathbf{U}_{(j,t)} = \max_j \mathbf{U}_{(j,t)} \\ 0 & \text{otherwise} \end{cases}$$

State labels from likelihood

$$\log \mathbf{U}_{(j,t)} = \log P(\tilde{\mathbf{x}}_t | \omega_{j_t})$$

Likelihood from uniform transition prob. (Naïve Bayes)



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 1: Decoding a Sequence

- Smoothing classification with transition probabilities

- For t -th frame

- Calculate the log likelihood $\log \mathbf{U}_{(j,t)} = \log P(\tilde{\mathbf{x}}_t | \omega_{j_t} = j)$

- Calculate the prior

- For this we need to know $\tilde{\mathbf{U}}_{(j,t-1)}$

$$P_{\mathbf{U}_{(j,t)}} = P(\omega_{j_t} | \omega_{j_{t-1}}) = P_{tr}^\top \tilde{\mathbf{U}}_{(:,t-1)} \quad \text{Keeps } j\text{-th column of } P_{tr}^\top$$

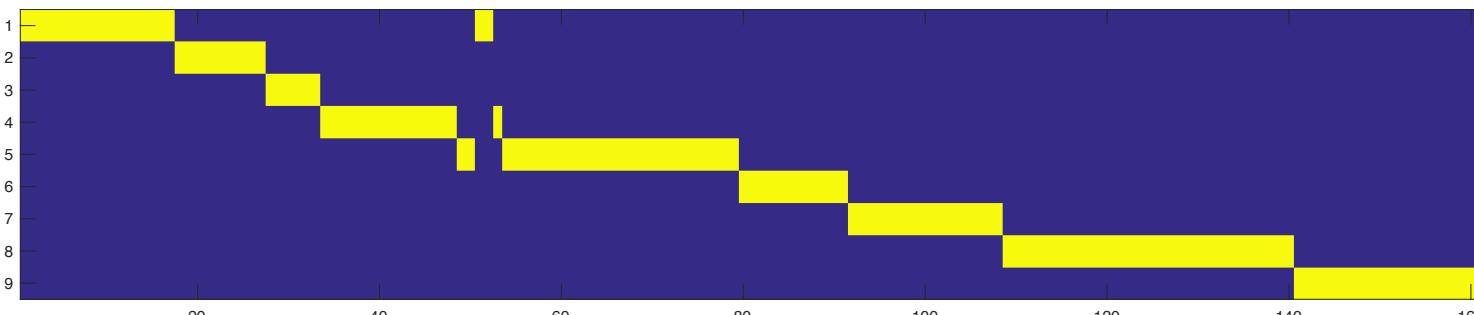
where j' is the label of the $(t-1)$ -th state

- Calculate the posterior $P_{\mathbf{U}_{(:,t)}} \odot \mathbf{U}_{(:,t)}$

- Find the final state $\tilde{\mathbf{U}}_{(j,t)} = \begin{cases} 1 & \text{if } P_{\mathbf{U}_{(j,t)}} \odot \mathbf{U}_{(j,t)} = \max_j P_{\mathbf{U}_{(j,t)}} \odot \mathbf{U}_{(j,t)} \\ 0 & \text{otherwise} \end{cases}$

- Do you like this approach?

- If you're wrong in the earlier frames, there's no way to fix it



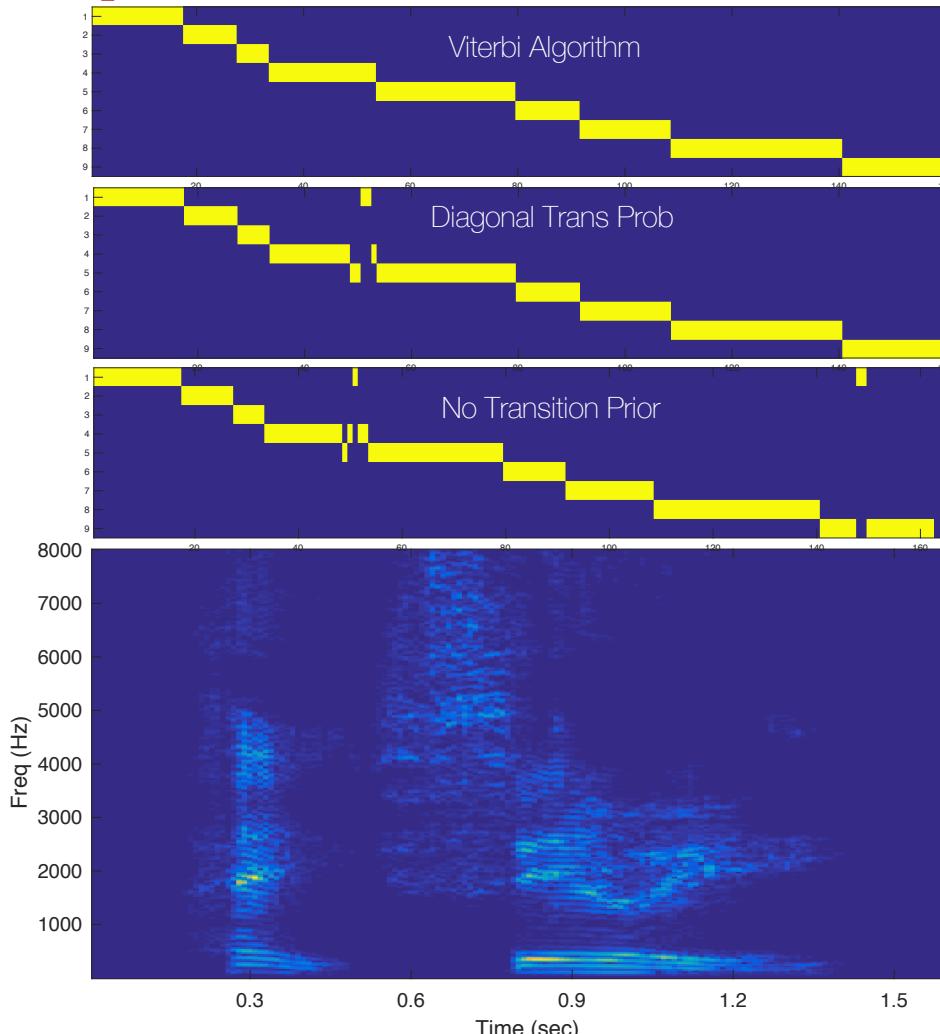
INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 1: Decoding a Sequence

- The Viterbi algorithm

- With the Viterbi algorithm we find the **best path**
 - Calculate the posterior probabilities and back track



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

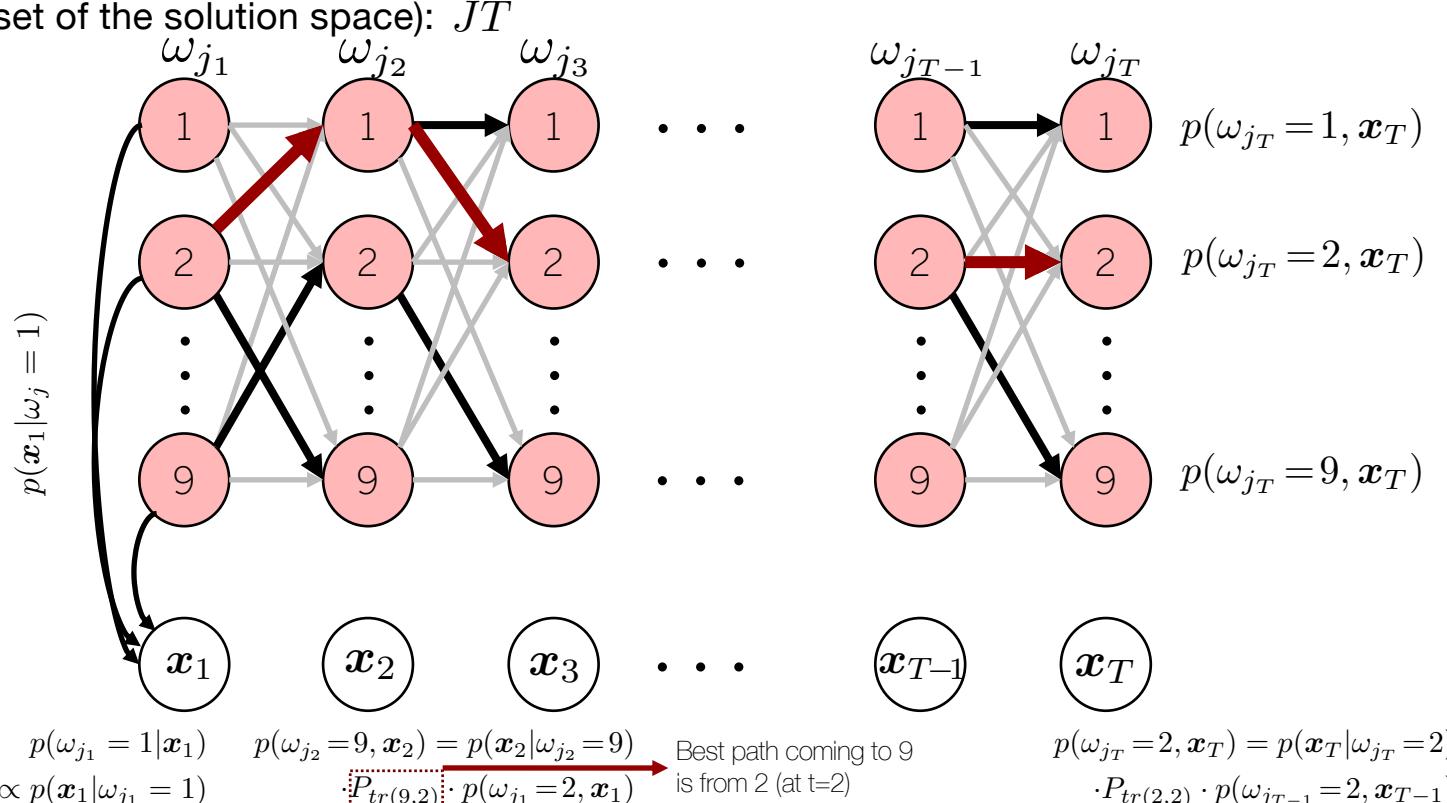
Case 1: Decoding a Sequence

- The Viterbi algorithm

- A dynamic programming algorithm

- The entire solution space: J^T

- Our candidates (subset of the solution space): JT



INDIANA UNIVERSITY

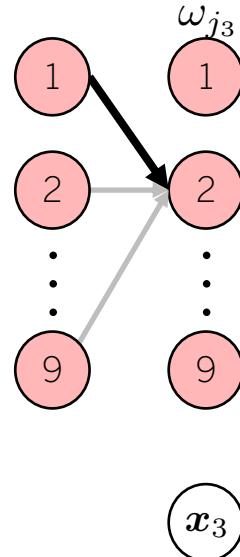
SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 1: Decoding a Sequence

- The Viterbi algorithm

- The Viterbi algorithm:
- Prepare the log-likelihoods $\log U_{(j,t)} = \log P(\tilde{x}_t | \omega_{jt})$
- Calculate the posterior (joint) probability from the left (earlier frames) to the right
 - For j -th state at t -th frame there are J different possible paths from the previous frame
 - The Viterbi algorithm chooses and keeps the best path from $t-1$ to t
$$p(\omega_{jt}, \mathbf{x}_t) \propto \max_{j_{t-1}} p(\mathbf{x}_t | \omega_{jt}) \cdot p(\omega_{jt} | \omega_{j_{t-1}}) \cdot p(\omega_{j_{t-1}}, \mathbf{x}_{t-1})$$
$$\log p(\omega_{jt}, \mathbf{x}_t) = \max_{j_{t-1}} \log p(\mathbf{x}_t | \omega_{jt}) + \log p(\omega_{jt} | \omega_{j_{t-1}}) + \log p(\omega_{j_{t-1}}, \mathbf{x}_{t-1}) + const.$$
- Backtracking
 - The final J states at $t=T$ have their final posterior probabilities
 - Pick the best $p(\omega_{j_T}, \mathbf{x}_T)$, e.g. $\arg \max_{j_T} p(\omega_{j_T}, \mathbf{x}_T) = 9$
 - What was $\omega_{j_{T-1}}$ that led to it?
 - What was $\omega_{j_{T-2}}$ that led to the best?
 - ...
- This procedure is based on a recursive assumption

$$\log p(\omega_{jt}, \mathbf{x}_t) = \log p(\mathbf{x}_t | \omega_{jt}) + \max_{j_{t-1}} \{ \log p(\omega_{j_{t-1}}, \mathbf{x}_t) + \log p(\omega_{jt} | \omega_{j_{t-1}}) \}$$



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 1: Decoding a Sequence

- The Viterbi algorithm

- You have a model, maybe a naïve Bayes
 - Or, any model that assumes a set of latent variables (classes)
 - The model also assumes each class has a unique generative model and observations are from them
 - Without any temporal structure this is a classification problem
- With the n-th order Markov model on the states

$$P(\omega_{j_t} | \omega_{j_{t-1}}, \omega_{j_{t-2}}, \dots, \omega_{j_1}) = P(\omega_{j_t} | \omega_{j_{t-1}})$$

- This turns into a decoding problem
- It's still a classification problem
 - But you need to smooth your labels based on the transition probabilities
- The Viterbi algorithm does this labeling job in two paths
 - Better than naïve forward-path-only smoothing
- The decoding problem assumes
 - You know the LV-specific generative models
 - You know the transition probabilities
 - **What if you don't?**



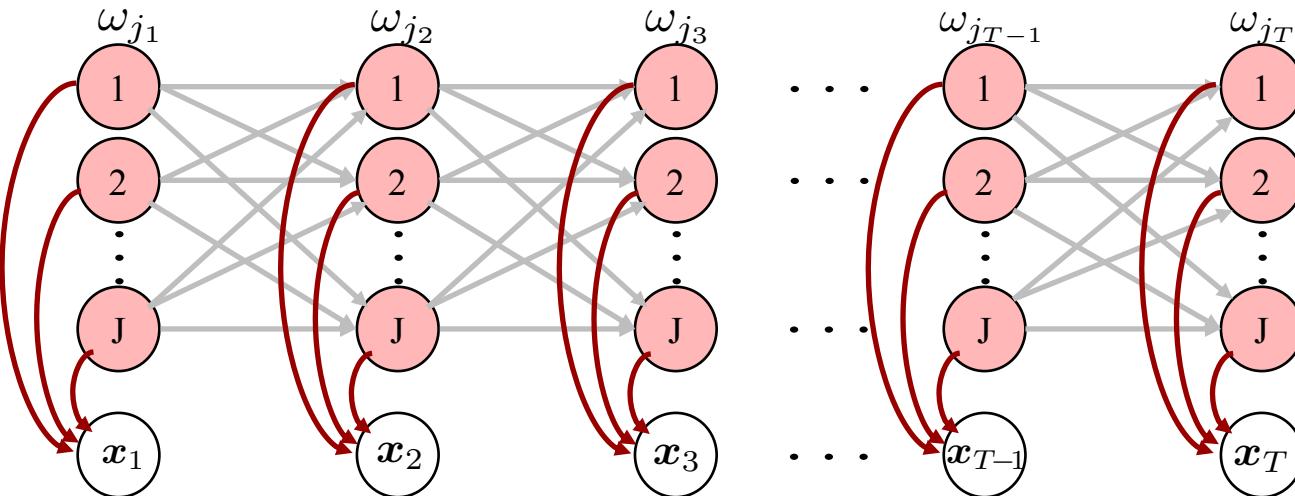
INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 2: Learning from a Sequence

- Hidden Markov Models (HMM)

- Let me properly define an HMM first



- The observations: $x_t \in \mathbb{R}^{D \times 1}$
- The hidden states: $\omega_{j_t} \leftarrow$ Decoding (Viterbi) procedure estimates this (given the following)
- Transition probabilities: $p(\omega_{j_t} | \omega_{j_{t-1}}) \in \mathbb{R}^{J \times J} \rightarrow \leftarrow$ HMM learning estimates this
- Emission probabilities: $p(x_t | \omega_{j_t}) \in \mathbb{R}^{D \times J} \curvearrowleft \leftarrow$ HMM learning estimates this
- Initial priors: $p(\omega_{j_1}) \in \mathbb{R}^{J \times 1} \leftarrow$ HMM learning estimates this



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 2: Learning from a Sequence

- The Viterbi reestimation algorithm

- Initialize $p(\omega_{j_t} | \omega_{j_{t-1}})$ with random numbers

- Prepare the log-likelihoods $\log U_{(j,t)} = \log P(\tilde{x}_t | \omega_{j_t})$

- Calculate the posterior probability using the local best paths

$$\log p(\omega_{j_t}, \mathbf{x}_t) = \max_{j_{t-1}} \log p(\mathbf{x}_t | \omega_{j_t}) + \log p(\omega_{j_t} | \omega_{j_{t-1}}) + \log p(\omega_{j_{t-1}}, \mathbf{x}_{t-1}) + \text{const.}$$

- Backtracking

- Estimate the best path given the data and the parameter estimation:

$$\Omega_i = \{\omega_{j_1} = 1, \omega_{j_2} = 1, \dots, \omega_{j_{t-1}} = 3, \omega_{j_t} = 4, \omega_{j_{t+1}} = 4, \dots, \omega_{j_T} = 9\}$$

- Then, it's easy to count these values

- $n_{i|j}$: number of transitions from state j to i
 - $n_{|j}$: number of transitions originated from state j

- Finally $p(\omega_{j_t} | \omega_{j_{t-1}}) \leftarrow \frac{n_{i|j}}{n_{|j}}$



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

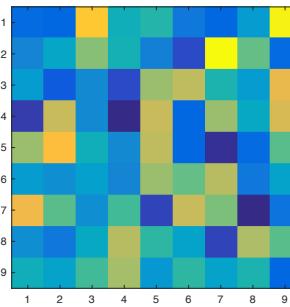
Case 2: Learning from a Sequence

- The Viterbi reestimation algorithm

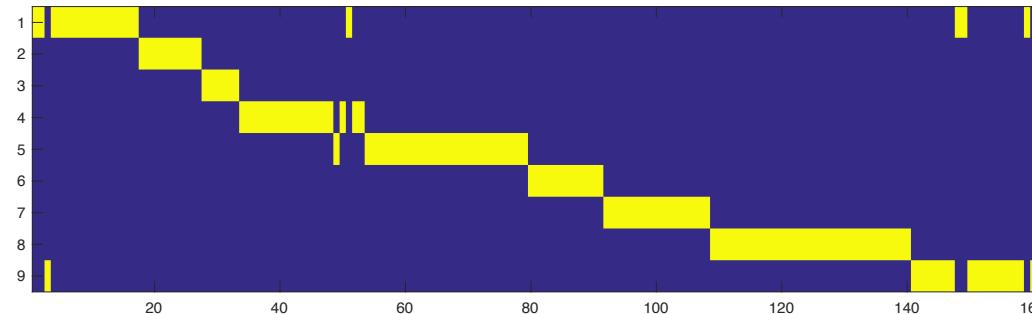
- Demo

$$p(\omega_{j_t} | \omega_{j_{t-1}})$$

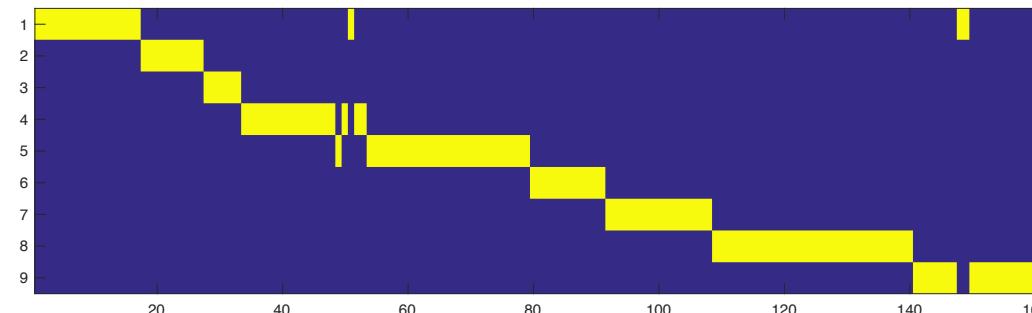
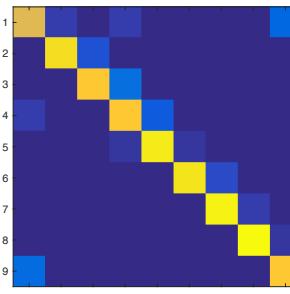
Iteration 1



Estimated States



Iteration 2



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

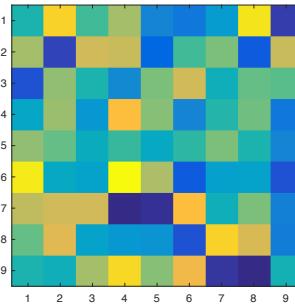
Case 2: Learning from a Sequence

- The Viterbi reestimation algorithm

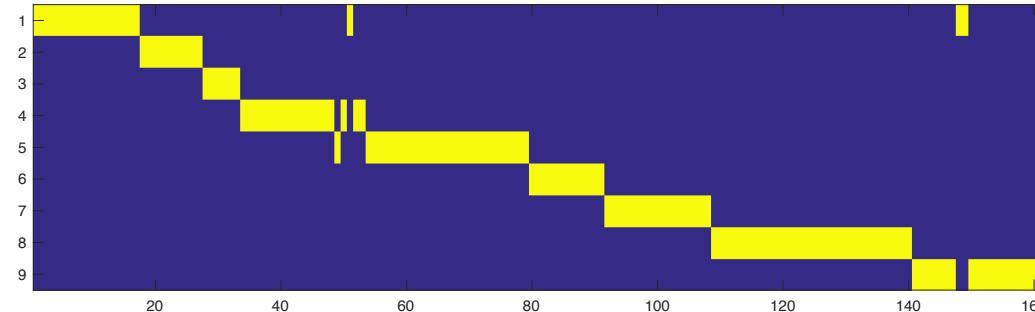
- Prior for the diagonals (erase off-diagonals at every iteration)

$$p(\omega_{j_t} | \omega_{j_{t-1}})$$

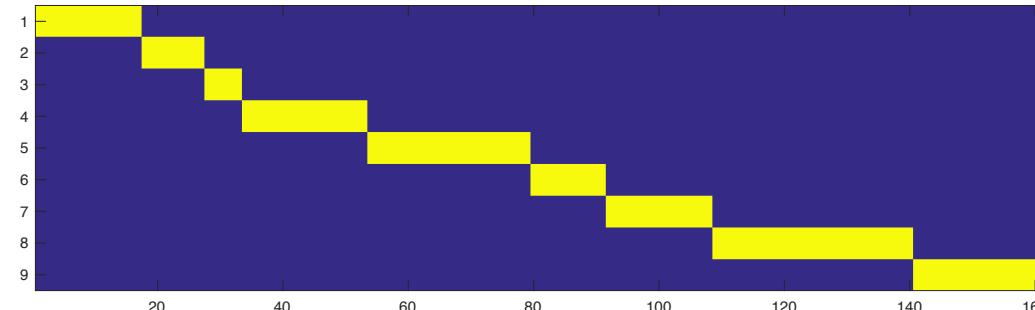
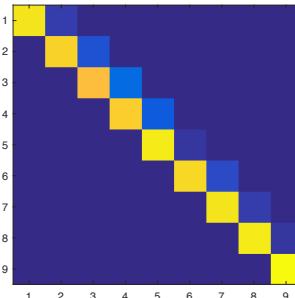
Iteration 1



Estimated States



Iteration 2



INDIANA UNIVERSITY

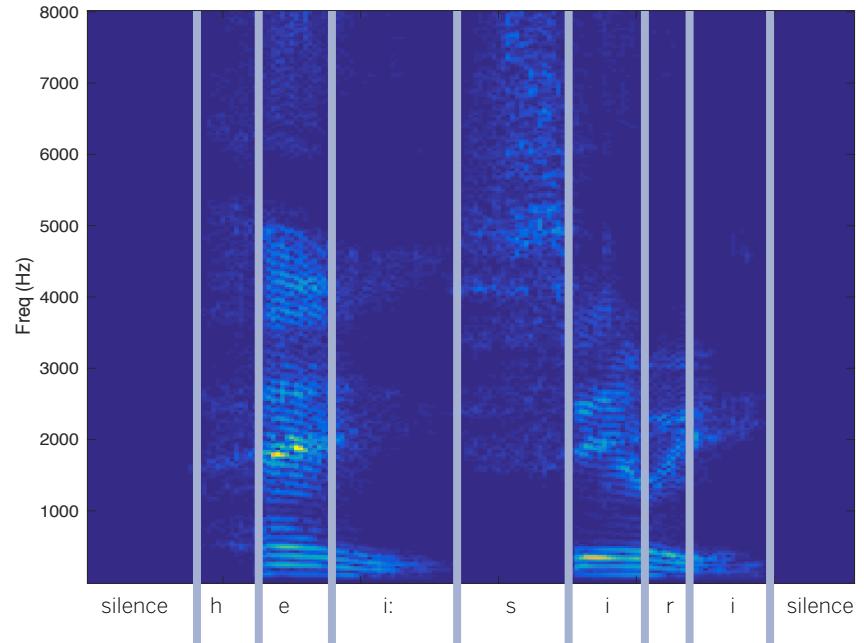
SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 2: Learning from a Sequence

- The Viterbi reestimation algorithm

- Why was this so easy?
 - We saw a convergence after only one iteration
 - B/C we started from a good set of emission probabilities using training data
 - If you remember, I manually cluster the MFCC frames to get the cluster-specific distribution
- But HMM learning can be seen as a clustering problem (or a mixture of distributions)
 - So, we must be able to learn the other parameters, too
 - I mean the emission probabilities

$$p(\mathbf{x}_t | \omega_{j_t}) \in \mathbb{R}^{D \times J}$$



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 2: Learning from a Sequence

- The Viterbi reestimation algorithm (full)

- Initialize $p(\omega_{j_t} | \omega_{j_{t-1}})$ with random numbers
- Prepare the log-likelihoods $\log U_{(j,t)} = \log P(\tilde{x}_t | \omega_{j_t})$
- Calculate the posterior probability using the best paths
$$\log p(\omega_{j_t}, \mathbf{x}_t) = \max_{j_{t-1}} \log p(\mathbf{x}_t | \omega_{j_t}) + \log p(\omega_{j_t} | \omega_{j_{t-1}}) + \log p(\omega_{j_{t-1}}, \mathbf{x}_{t-1}) + const.$$
- Backtracking
 - Estimate the best path given the data and the parameter estimation:
$$\Omega_i = \{\omega_{j_1} = 1, \omega_{j_2} = 1, \dots, \omega_{j_{t-1}} = 3, \omega_{j_t} = 4, \omega_{j_{t+1}} = 4, \dots, \omega_{j_T} = 9\}$$
- Then, it's easy to count these values
 - $n_{i|j}$: number of transitions from state j to i
 - $n_{|j}$: number of transitions originated from state j
 - $n_i|$: number of transitions terminated at state i
- Finally $p(\omega_{j_t} | \omega_{j_{t-1}}) \leftarrow \frac{n_{i|j}}{n_{|j}}$
- For the observations associated with state j , do your parameter estimation
 - e.g. Maximum likelihood estimation
 - In my case the emission probability is $p(\mathbf{x}_t | \omega_{j_t}) = \mathcal{N}(\mathbf{x}_t; \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)$
 - Therefore,
$$\boldsymbol{\mu}_i = \frac{1}{n_i|} \sum_{t=1}^T \mathcal{I}(\omega_{j_t} = i) \mathbf{x}_t$$

$$\mathcal{I}(x) = \begin{cases} 1 & \text{if } x \text{ is true} \\ 0 & \text{otherwise} \end{cases}$$

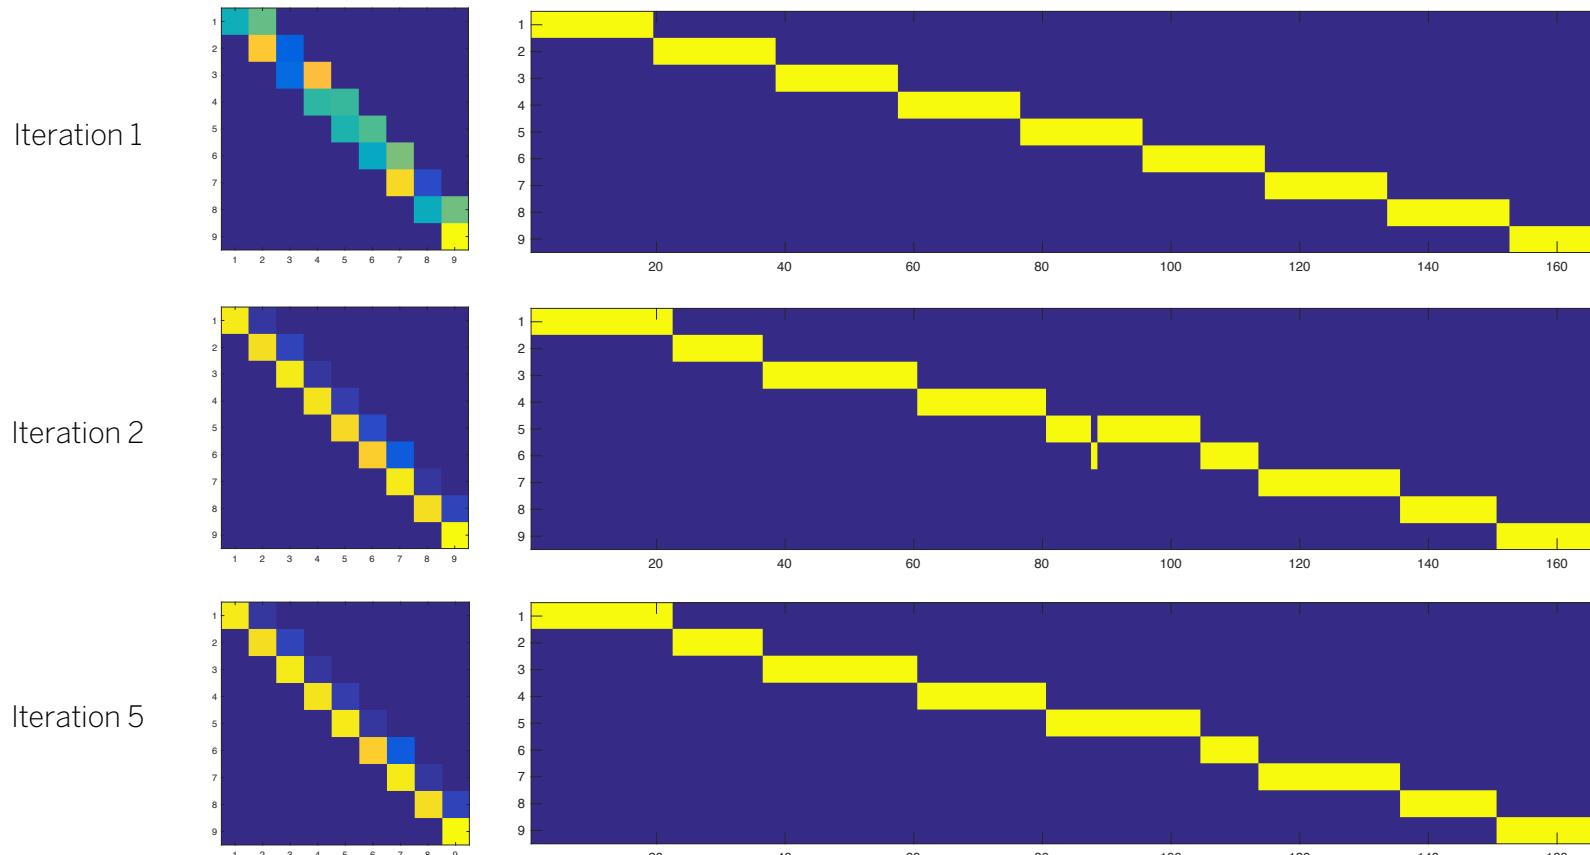


INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 2: Learning from a Sequence

- The Viterbi reestimation algorithm (full)
 - This time the Gaussian parameters are updated as well



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 3: Evaluating a Sequence

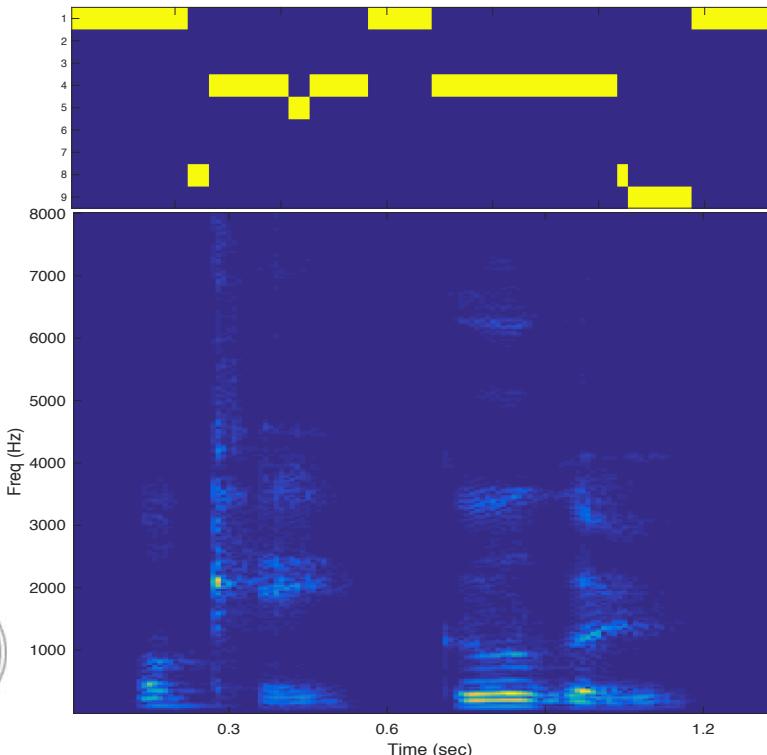
- Recognition using the best path

- Now you know how to find the best path for a test signal
 - And it comes with a posterior probability of the state transitions

$$p(\Omega_i)p(\mathbf{X}|\Omega_i) \quad \Omega_i : i\text{-th (best) path out of } J^T \text{ possible paths}$$

- Another test signal
 - Log of the posterior prob of the best path:
-13469
 - Compare this with that of the test signal that says “Hey, Siri”:
-3230

- The best path found is not good enough to model the new signal
 - The new test signal might be far from the training signal
- Can we be more careful?



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 3: Evaluating a Sequence

- Recognition using all paths

- The forward pass:
- Prepare the log-likelihoods $\log U_{(j,t)} = \log P(\tilde{\mathbf{x}}_t | \omega_{jt})$
- ~~Calculate the joint probability from the left (earlier frames) to the right~~
 - ~~For j -th state at t -th frame there are J different possible paths from the previous frame~~
 - ~~The Viterbi algorithm chooses and keeps the best path from $t-1$ to t~~
 - $\log p(\omega_{jt}, \mathbf{x}_t) = \max_{j_{t-1}} \log p(\mathbf{x}_t | \omega_{jt}) + \log p(\omega_{jt} | \omega_{j_{t-1}}) + \log p(\omega_{j_{t-1}}, \mathbf{x}_{t-1}) + \text{const.}$

- The forward algorithm calculates from all paths and sum them up

$$\alpha_{j,1} = P(\omega_{j_1})p(\mathbf{x}_1 | \omega_{j_1}) = p(\mathbf{x}_1, \omega_{j_1})$$

$$\begin{aligned}\alpha_{j,2} &= p(\mathbf{x}_1, \mathbf{x}_2, \omega_{j_2}) = \sum_{\omega_{j_1}} p(\mathbf{x}_1, \mathbf{x}_2, \omega_{j_2}, \omega_{j_1}) = p(\mathbf{x}_2 | \omega_{j_2}) \sum_{\omega_{j_1}} p(\mathbf{x}_1, \omega_{j_1})p(\omega_{j_2} | \omega_{j_1}) \\ &= p(\mathbf{x}_2 | \omega_{j_2}) \sum_{\omega_{j_1}} \alpha_{j,1} p(\omega_{j_2} | \omega_{j_1})\end{aligned}$$

$$\begin{aligned}\alpha_{j,3} &= p(\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3, \omega_{j_3}) = \sum_{\omega_{j_2}} p(\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3, \omega_{j_3}, \omega_{j_2}) = p(\mathbf{x}_3 | \omega_{j_3}) \sum_{\omega_{j_2}} p(\mathbf{x}_1, \mathbf{x}_2, \omega_{j_2})p(\omega_{j_3} | \omega_{j_2}) \\ &= p(\mathbf{x}_3 | \omega_{j_3}) \sum_{\omega_{j_2}} \alpha_{j,2} p(\omega_{j_3} | \omega_{j_2})\end{aligned}$$

- Where am I going? $\alpha_{j,t} = p(\mathbf{x}_t | \omega_{j_t}) \sum_{\omega_{j_{t-1}}} \alpha_{j,t-1} p(\omega_{j_t} | \omega_{j_{t-1}})$



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 3: Evaluating a Sequence

- Recognition using all paths

- The forward pass

$$\alpha_{j,t} = p(\mathbf{x}_t | \omega_{j_t}) \sum_{\omega_{j_{t-1}}} \alpha_{j,t-1} p(\omega_{j_t} | \omega_{j_{t-1}})$$

- For your information..

- If you replace the summation with maximum

$$\alpha_{j,t} = p(\mathbf{x}_t | \omega_{j_t}) \max_{j_{t-1}} \alpha_{j,t-1} p(\omega_{j_t} | \omega_{j_{t-1}})$$

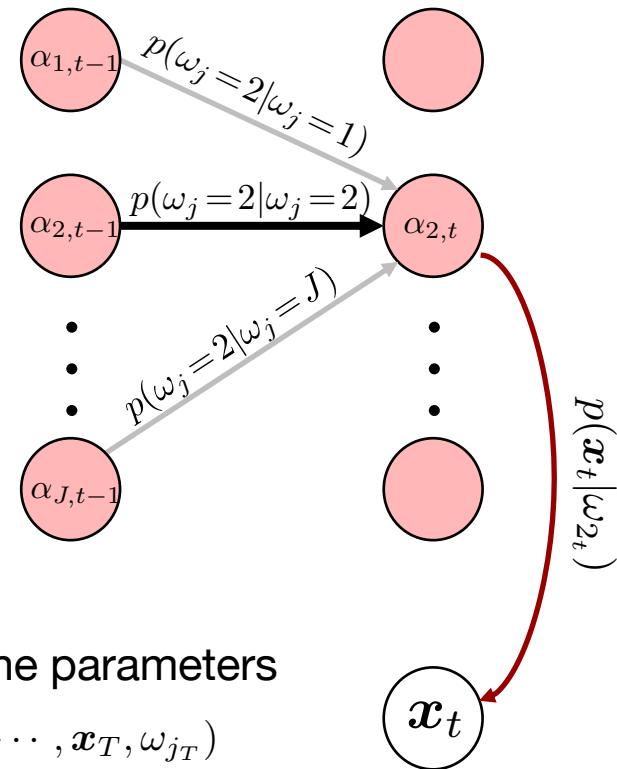
- It corresponds to the forward pass of the Viterbi algorithm

- In the end, we get the probability of observing the data given the parameters

$$\alpha_{j,T} = p(\mathbf{x}_T | \omega_{j_T}) \sum_{\omega_{j_{T-1}}} \alpha_{j,T-1} p(\omega_{j_T} | \omega_{j_{T-1}}) = p(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T, \omega_{j_T})$$

$$p(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T) = \sum_j \alpha_{j,T} \quad p(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T; \Theta) = \sum_j \alpha_{j,T}$$

- If $\sum_j \alpha_{j,T}$ is high, the model fits the data better



Case 2: Learning from a Sequence (Revisited)

- The Baum-Welch algorithm

- Now we have a soft version of the forward pass

- Let's use this for the learning task

- The backward pass

$$\beta_{j,T} = 1$$

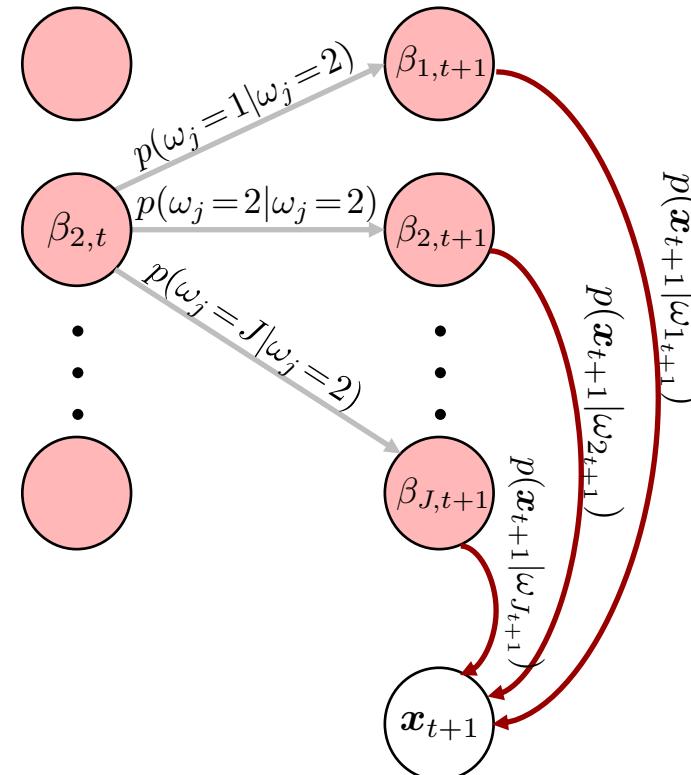
$$\beta_{j,t} = p(\mathbf{x}_{t+1}, \dots, \mathbf{x}_T | \omega_{j_t})$$

$$= \sum_{\omega_{j_{t+1}}} p(\mathbf{x}_{t+1}, \dots, \mathbf{x}_T, \omega_{j_{t+1}} | \omega_{j_t})$$

$$= \sum_{\omega_{j_{t+1}}} p(\mathbf{x}_{t+1} | \omega_{j_{t+1}}) p(\mathbf{x}_{t+2}, \dots, \mathbf{x}_T | \omega_{j_{t+1}}) p(\omega_{j_{t+1}} | \omega_{j_t})$$

$$= \sum_{\omega_{j_{t+1}}} p(\mathbf{x}_{t+1} | \omega_{j_{t+1}}) \beta_{j_{t+1}} p(\omega_{j_{t+1}} | \omega_{j_t})$$

- Let's combine the passes

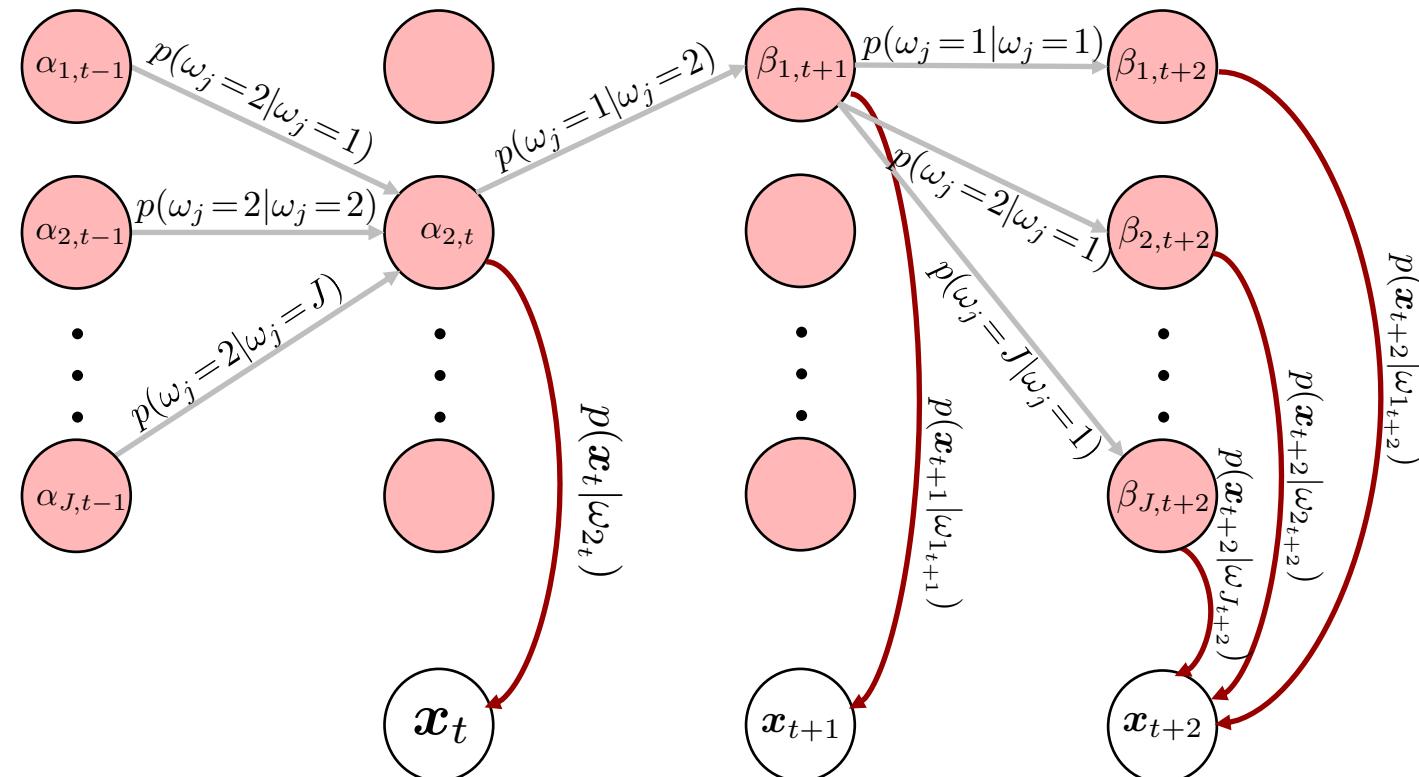


Case 2: Learning from a Sequence (Revisited)

- The Baum-Welch algorithm

○

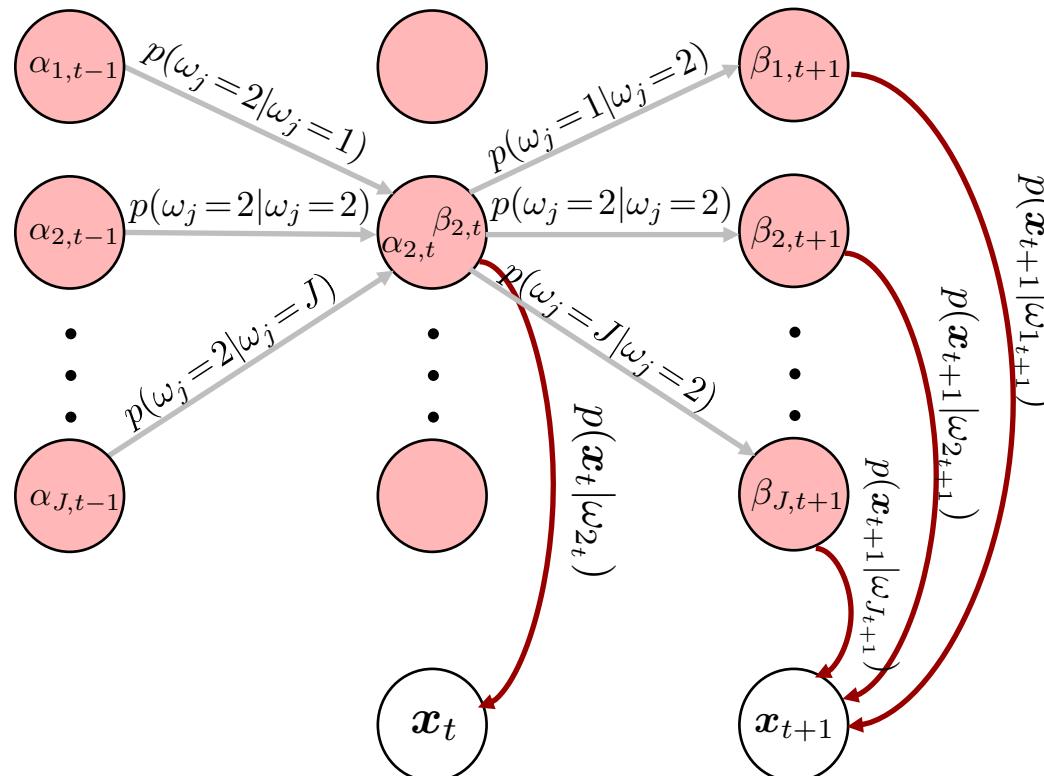
$$p(\mathbf{x}_1, \dots, \mathbf{x}_t, \omega_{2_t}, \mathbf{x}_{t+1}, \omega_{1_{t+1}}, \mathbf{x}_{t+2}, \dots, \mathbf{x}_T) = \alpha_{2,t} p(\omega_j=1|\omega_j=2) p(\mathbf{x}_{t+1}|\omega_{1_{t+1}}) \beta_{1,t+1}$$



Case 2: Learning from a Sequence (Revisited)

- The Baum-Welch algorithm

- $p(\mathbf{x}_1, \dots, \mathbf{x}_T, \omega_{2,t}) = \alpha_{2,t} \beta_{2,t}$



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 2: Learning from a Sequence (Revisited)

- The Baum-Welch algorithm

$$p(\mathbf{x}_1, \dots, \mathbf{x}_T, \omega_{2_t}) = \alpha_{2,t} \beta_{2,t}$$

$$p(\mathbf{x}_1, \dots, \mathbf{x}_t, \omega_{2_t}, \mathbf{x}_{t+1}, \omega_{1_{t+1}}, \mathbf{x}_{t+2}, \dots, \mathbf{x}_T) = \alpha_{2,t} p(\omega_j = 1 | \omega_j = 2) p(\mathbf{x}_{t+1} | \omega_{1_{t+1}}) \beta_{1,t+1}$$

- Transition matrix estimation

- Get the pseudocount

$$p(\omega_1 | \omega_2) \leftarrow \frac{\sum_t p(\mathbf{x}_1, \dots, \mathbf{x}_T, \omega_{2_t}, \omega_{1_{t+1}})}{\sum_t p(\mathbf{x}_1, \dots, \mathbf{x}_T, \omega_{2_t})} = \frac{\sum_t \alpha_{2,t} p(\omega_j = 1 | \omega_j = 2) p(\mathbf{x}_{t+1} | \omega_{1_{t+1}}) \beta_{1,t+1}}{\sum_t \alpha_{2,t} \beta_{2,t}}$$

- Emission probability

- Use the posterior prob $p(\omega_{j_t} | \mathbf{x}_1 \dots \mathbf{x}_T) \propto \gamma_{j,t} = \alpha_{j,t} \beta_{j,t}$

- Remember $p(\mathbf{x}_t | \omega_{j_t}) = \mathcal{N}(\mathbf{x}_t; \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)$

$$\boldsymbol{\mu}_j = \frac{\sum_{t=1}^T \gamma_{j,t} \mathbf{x}_t}{\sum_{t=1}^T \gamma_{j,t}}$$

$$\boldsymbol{\mu}_{j'} = \frac{1}{n_{i'}} \sum_{t=1}^T \mathcal{I}(\omega_{j_t} = j') \mathbf{x}_t$$

$$\mathcal{I}(x) = \begin{cases} 1 & \text{if } x \text{ is true} \\ 0 & \text{otherwise} \end{cases}$$

- EM with a more complex E-step!



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Case 2: Learning from a Sequence (Revisited)

- The Baum-Welch algorithm

- What happens if my emission probabilities are not following Gaussian?
 - There's no guarantee
 - Is there a flexible probabilistic distribution that can model ANY data?
 - Gaussian mixture models
- Eventually, a full HMM will be
 - A mixture of GMM
 - Each state has a corresponding GMM as its emission probability
 - You do another EM for the parameter estimation
 - Instead of the MLE step



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Recap

- Three problems in Hidden Markov Models (HMM)
 - Decoding: given the observation sequence and model parameters, find out the best sequence of hidden states
 - I said it's classification by seeing each hidden state as a class, but in practice it's not
 - Evaluation: calculate the probability of observing the data given the model parameters
 - An HMM model per class is a more common classification setting
 - Learning: learning the model parameters from data
 - Corresponds to training
- Decoding
 - Viterbi algorithm
 - More efficient than an exhaustive search $O(J^T)$; uses the best path out of $O(JT)$; backtracking
- Evaluation
 - Using the best path
 - Do Viterbi and pick up the probability of the best path
 - Forward algorithm
 - Soften up the decision. Gather up all the other probabilities coming from the suboptimal paths.
- Learning
 - The Viterbi reestimation
 - Do Viterbi and count the number of frames per state. Use this to do the clustering to reestimate the parameters.
 - The Baum-Welch algorithm
 - Using the forward pass and backward pass, calculate the probability of having j -th state given the entire observation sequence. This is your E-step.



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING

Reading

- Chapter 9
- L. R. Rabiner, "A tutorial on hidden Markov models and selected applications in speech recognition," in *Proceedings of the IEEE*, vol. 77, no. 2, pp. 257-286, Feb 1989.
doi: 10.1109/5.18626



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING



Thank You!



INDIANA UNIVERSITY

SCHOOL OF INFORMATICS, COMPUTING, AND ENGINEERING